



✓ **Congratulations! You passed!**

TO PASS 80% or higher

Keep Learning

GRADE
100%

Practical aspects of deep learning

LATEST SUBMISSION GRADE

100%

1. If you have 10,000,000 examples, how would you split the train/dev/test set?

1 / 1 point

- ☒ 98% train . 1% dev . 1% test
- ☐ 60% train . 20% dev . 20% test
- ☐ 33% train . 33% dev . 33% test

✓ Correct

2. The dev and test set should:

1 / 1 point

- ☒ Come from the same distribution
- ☐ Come from different distributions
- ☐ Be identical to each other (same (x,y) pairs)
- ☐ Have the same number of examples

✓ Correct

3. If your Neural Network model seems to have high variance, what of the following would be promising things to try?

1 / 1 point

- ☒ Get more training data

✓ Correct

- ☐ Increase the number of units in each hidden layer
- ☐ Make the Neural Network deeper
- ☒ Add regularization

✓ Correct

- ☐ Get more test data

4. You are working on an automated check-out kiosk for a supermarket, and are building a classifier for apples, bananas and oranges. Suppose your classifier obtains a training set error of 0.5%, and a dev set error of 7%. Which of the following are promising things to try to improve your classifier? (Check all that apply.)

1 / 1 point

- ☒ Increase the regularization parameter lambda

✓ Correct

- ☐ Decrease the regularization parameter lambda
- ☒ Get more training data

✓ Correct

✓ Correct

☐ Use a bigger neural network

5. What is weight decay?

1 / 1 point

- ☐ A technique to avoid vanishing gradient by imposing a ceiling on the values of the weights.
- ☒ A regularization technique (such as L2 regularization) that results in gradient descent shrinking the weights on every iteration.
- ☐ The process of gradually decreasing the learning rate during training.
- ☐ Gradual corruption of the weights in the neural network if it is trained on noisy data.

✓ Correct

6. What happens when you increase the regularization hyperparameter lambda?

1 / 1 point

- ☒ Weights are pushed toward becoming smaller (closer to 0)
- ☐ Weights are pushed toward becoming bigger (further from 0)
- ☐ Doubling lambda should roughly result in doubling the weights
- ☐ Gradient descent taking bigger steps with each iteration (proportional to lambda)

✓ Correct

7. With the inverted dropout technique, at test time:

1 / 1 point

- ☐ You apply dropout (randomly eliminating units) and do not keep the 1/keep_prob factor in the calculations used in training
- ☐ You apply dropout (randomly eliminating units) but keep the 1/keep_prob factor in the calculations used in training.
- ☐ You do not apply dropout (do not randomly eliminate units), but keep the 1/keep_prob factor in the calculations used in training.
- ☒ You do not apply dropout (do not randomly eliminate units) and do not keep the 1/keep_prob factor in the calculations used in training

✓ Correct

8. Increasing the parameter keep_prob from (say) 0.5 to 0.6 will likely cause the following: (Check the two that apply)

1 / 1 point

- ☐ Increasing the regularization effect
- ☒ Reducing the regularization effect

✓ Correct

- ☐ Causing the neural network to end up with a higher training set error
- ☒ Causing the neural network to end up with a lower training set error

✓ Correct

9. Which of these techniques are useful for reducing variance (reducing overfitting)? (Check all that apply.)

1 / 1 point

- ☒ L2 regularization

✓ Correct

- ☒ Data augmentation

✓ Correct

- ☒ Dropout

✓ Correct

- ☐ Exploding gradient
- ☐ Gradient Checking
- ☐ Vanishing gradient
- ☐ Xavier initialization

10. Why do we normalize the inputs x ?

1 / 1 point

- ☒ It makes the cost function faster to optimize
- ☐ It makes it easier to visualize the data
- ☐ It makes the parameter initialization faster
- ☐ Normalization is another word for regularization--It helps to reduce variance

✓ Correct